

# Experiment 4: Query Evaluation Plan Analysis

## Prerequisites

Knowledge of SQL queries, JOINs, and indexing concepts

---

## Lab Objective

To analyze how SQL queries are executed internally by the database management system using query execution plans and understand the cost and performance implications of different query strategies.

---

## Learning Outcomes

By completing this experiment, students will be able to:

- Generate and interpret query execution plans.
  - Understand how the database optimizer chooses execution strategies.
  - Identify table scans, index scans, and join methods.
  - Compare performance before and after indexing.
  - Analyze query cost and execution time.
- 

## Introduction

When a SQL query is executed, the database does not directly run it as written. Instead, it generates an **execution plan** that determines how data will be accessed, joined, filtered, and sorted.

Query evaluation plans help developers understand:

- How tables are accessed
- Whether indexes are used
- Which join algorithms are selected
- Estimated cost and execution time

This experiment focuses on analyzing and interpreting execution plans to improve database performance.

---

## Key Concepts Covered

- **Query Optimizer:** Component that determines the most efficient way to execute a query.
  - **Execution Plan:** Step-by-step strategy used to retrieve data.
  - **Cost Estimation:** Estimated resource usage (CPU, I/O).
  - **Table Scan vs Index Scan:** Methods of data retrieval.
  - **Join Algorithms:** Nested Loop Join, Hash Join, Merge Join.
  - **Execution Time Analysis:** Comparing performance metrics.
- 

## Experiment Steps

---

### 1. Generate Execution Plan

Use the database's EXPLAIN feature to view query execution plans.

#### Example:

```
EXPLAIN SELECT *
  FROM Students s
  JOIN Enrollments e ON s.student_id = e.student_id
 WHERE s.age > 20;
```

Observe:

- Access type (ALL, index, range, etc.)
  - Rows examined
  - Possible keys and selected keys
  - Join order
- 

## 2. Analyze Table Access Methods

Identify whether:

- Full table scans are being used.
- Indexes are being utilized.

Example observation:

- If access type = ALL → Full table scan
- If access type = index or range → Index usage

Students should record differences in row count and cost.

---

## 3. Study Join Operations

Observe different join algorithms used by the optimizer.

For large datasets, the system may choose:

- Nested Loop Join
- Hash Join
- Merge Join

Example:

```
EXPLAIN SELECT s.name, c.course_name  
FROM Students s  
JOIN Enrollments e ON s.student_id = e.student_id  
JOIN Courses c ON e.course_id = c.course_id;
```

Analyze:

- Join order
  - Join type
  - Estimated cost
- 

## 4. Compare Before and After Indexing

### Step 1 – Without Index

Run EXPLAIN on a query filtering by a non-indexed column.

### Step 2 – Create Index

```
CREATE INDEX idx_age ON Students(age);
```

### Step 3 – Re-evaluate Execution Plan

```
EXPLAIN SELECT * FROM Students WHERE age > 20;
```

Compare:

- Access type
  - Number of rows scanned
  - Query cost
  - Execution time
- 

## 5. Performance Comparison

Modify queries and observe changes in execution plan.

Examples:

- Add filtering conditions.
- Reorder JOIN clauses.
- Reduce selected columns instead of using SELECT \*.

Analyze how small changes affect:

- Estimated cost
  - Rows processed
  - Join method selection
- 

## How It Works

- The query optimizer evaluates multiple execution strategies.
  - It selects the plan with the lowest estimated cost.
  - Execution plans show how data is retrieved and processed internally.
  - Indexes significantly influence execution efficiency.
  - Understanding execution plans helps developers write optimized queries.
- 

## Submission Instructions

Students must submit:

- Screenshot of execution plan for a sample query.
  - Screenshot showing table scan or index usage.
  - Comparison of execution plan before and after indexing.
  - Brief analysis explaining performance differences observed.
- 

## Conclusion

This experiment enables students to understand the internal workings of query execution and optimization. The ability to interpret execution plans is a critical skill for database developers, database administrators, and performance engineers working with large-scale data systems.